



---

# *Journal of Statistical Software*

MMMMMM YYYY, Volume VV, Issue II.

*doi: 10.18637/jss.v000.i00*

---

## **cubble: An R Package for Organizing and Wrangling Multivariate Spatio-temporal Data**

**H. Sherry Zhang**  
Monash University

**Dianne Cook**  
Monash University

**Ursula Laa**  
University of Natural  
Resources and Life Sciences

**Nicolas Langrené**  
BNU-HKBU  
United International College

**Patricia Menéndez**  
Monash University

---

### **Abstract**

Multivariate spatio-temporal data refers to multiple measurements taken across space and time. For many analyses, spatial and time components can be separately studied: for example, to explore the temporal trend of one variable for a single spatial location, or to model the spatial distribution of one variable at a given time. However for some studies, it is important to analyze different aspects of the spatio-temporal data simultaneously, like for instance, temporal trends of multiple variables across locations. In order to facilitate the study of different portions or combinations of spatio-temporal data, we introduce a new class, **cubble**, with a suite of functions enabling easy slicing and dicing on different spatio-temporal components. The proposed **cubble** class ensures that all the components of the data are easy to access and manipulate while providing flexibility for data analysis. In addition, the **cubble** package facilitates visual and numerical explorations of the data while easing data wrangling and modelling. The **cubble** class and the tools implemented in the package are illustrated with different examples of Australian climate data.

*Keywords:* spatial, temporal, spatio temporal, R, environmental data, exploratory data analysis.

---

### **1. Introduction**

Spatio-temporal data has a spatial component referring to the location of each observation and a temporal component that is recorded at regular or irregular time intervals. It may also in-

## 2 **cubble**: An R Package for Organizing and Wrangling Multivariate Spatio-temporal Data

clude multiple variables measured at each spatial and temporal values. With spatio-temporal data, one can fix the time to explore the spatial features of the data, fix the spatial location/s to explore temporal aspects, or dynamically explore the space and time simultaneously.

In order to computationally explore the spatial, temporal and spatio-temporal faces of such data, the data needs to be stored and represented under a specific data object that allows the user to query, group and dissect all the data faces.

The Comprehensive R Archive Network (CRAN) task view SpatioTemporal (Pebesma and Bivand 2022) gathers information about R packages designed for spatio-temporal data and it has a section on *Representing data* that lists existing spatio-temporal data representations used in R. Among them, the **spacetime** package (Pebesma 2012) implements four S4 classes to handle spatio-temporal data with different spatio-temporal layouts (full grid, sparse grid, irregular, and trajectory). The **stars** package (Pebesma 2021) implements an S3 class built from dense arrays.

Still, these two implementations are not necessarily easy to work in analysis for analysts with a tidyverse mindset. In tidyverse, data are in tables and the tidy data concept (Wickham 2014) prescribes three principles on how data should be organised for easier analysis as 1) one observation a row, 2) one variable a column, and 3) one type of observation a table. The third principle of tidy data is particularly relevant for spatio-temporal data since spatial and temporal data are naturally observed at different units: the location and location at different times. While the tidyverse suite implements data wrangling and visualisation tools operated on a single table, there has not been many tools for handling relational data for spatio-temporal analysis. This motivates a new design to organise spatio-temporal data in a way that would make data wrangling, visualizing and analyzing easier.

This paper presents a new R package, **cubble**, which implements a new cubble class to organize spatial and temporal variables as two forms of a single data object so that they can be wrangled separately or combined, while being kept synchronized. Among the four spacetime layouts in Pebesma (2012), the **cubble** class can handle the full grid layout and the sparse grid layout. The software is available from the Comprehensive R Archive Network (CRAN) at <https://CRAN.R-project.org/package=cubble>.

The rest of the paper is organized as follows: Section 2 presents the main design and functionality of the **cubble** package. Section 3 explains how the **cubble** package deals with more advanced considerations, including data matching and how the package fits with existing static and interactive visualization tools. Moreover we also illustrate how the **cubble** package deals with spatio-temporal data transformations. Section 4 uses primarily Australian weather station data as examples to demonstrate the use of the package. An example of how the **cubble** package handles Network Common Data Form (NetCDF) data is also provided. Section 5 discuss the paper contributions and future directions.

## 2. The **cubble** package

### 2.1. The **cubble** object

Spatio-temporal data can encompass data with various spatial and temporal characteristics and different data require different structures for wrangling and analysis: climate weather stations record data at fixed point location but may suffer from potential temporal data

quality issue such as missing data for certain days. GPS data tracks unique point locations at different timestamps and is represented as trajectories. Satellite imageries capture snapshots of landscapes at selected time and is commonly structured as raster data. The spatio-temporal data cubble address are those collected at unique fixed locations while allowing for irregularity in the temporal dimension, such as the weather station data.

The cubble class is an S3 class built on tibble that allows the spatio-temporal data to be wrangled in two forms: a nested form (spatial cubble) and a long form (temporal cubble). It consists of two subclasses:

- a spatial cubble with class c("spatial\_cubble\_df", "cubble\_df")
- a temporal cubble with class c("temporal\_cubble\_df", "cubble\_df")

In a spatial cubble, spatial variables are organised as columns and temporal variables are nested within a specialised `ts` column. The cubble object printed below contains weather records in three airport stations (the creation of a cubble object will be explained in Section 2.2). This toy data is a subset of a larger data `climate_aus` collected from Global Historical Climatology Network Daily (GHCND). It records three airport stations located in Melbourne, Australia and includes spatial variables such as station ID, longitude, latitude, elevation, station name, World Meteorology Organisation ID. The dataset contains temporal variables including precipitation, maximum and minimum temperature, which can be read from the cubble header:

```
R> cb_nested
```

```
# cubble: key: id [3], index: date, nested form
# spatial: [144.8321, -37.98, 145.0964, -37.6655], Missing CRS!
# temporal: date [date], prcp [dbl], tmax [dbl], tmin [dbl]
  id      long   lat elev name          wmo_id ts
  <chr>    <dbl> <dbl> <dbl> <chr>        <dbl> <list>
1 ASN00086038 145. -37.7  78.4 essendon airport  95866 <tibble [10 x 4]>
2 ASN00086077 145. -38.0  12.1 moorabbin airport 94870 <tibble [10 x 4]>
3 ASN00086282 145. -37.7 113. melbourne airport 94866 <tibble [10 x 4]>
```

In a temporal cubble, temporal variables are expanded in the long form and spatial variables are stored as a data attribute:

```
R> cb_long
```

```
# cubble: key: id [3], index: date, long form
# temporal: 2020-01-01 -- 2020-01-10 [1D], no gaps
# spatial: long [dbl], lat [dbl], elev [dbl], name [chr], wmo_id [dbl]
  id      date      prcp  tmax tmin
  <chr>    <date>    <dbl> <dbl> <dbl>
1 ASN00086038 2020-01-01     0  26.8  11
2 ASN00086038 2020-01-02     0  26.3  12.2
3 ASN00086038 2020-01-03     0  34.5  12.7
```

## 4 **cubble**: An R Package for Organizing and Wrangling Multivariate Spatio-temporal Data

```
4 ASN00086038 2020-01-04      0  29.3  18.8
5 ASN00086038 2020-01-05     18  16.1  12.5
# i 25 more rows
```

### *The cubble attributes*

A cubble object inherits the attributes from tibble (and its subclasses): `class`, `row.names`, and `names`. Additionally, it has three specialised attributes: `key`, `index`, and `coords`. Readers who are familiar with the `key` and `index` attributes from the `tsibble` package would already understand the two arguments. In cubble, the `key` attribute identifies the row in the spatial cubble (given the use of `tidy::nest()` for nesting), and when combined with the `index` argument, it identifies the row in the temporal cubble. Currently, cubble only supports one variable as the key, and the accepted temporal classes for index includes the base R classes `Date`, `POSIXlt`, `POSIXct`, as well as `tsibble`'s `yearmonth`, `yearweek`, and `yearquarter` classes. The `coords` attribute represents an ordered pair of coordinates. It can be either an unprojected pair of longitude and latitude, or a projected easting and northing value. The `sf` package is used under the hood to calculate the bounding box, displayed in the header of a spatial cubble.

The temporal cubble has a special attribute called `spatial` to store the spatial variables. Shortcut functions are available to extract attributes, for example, `spatial()` for extracting spatial variables from the temporal cubble:

```
R> spatial(cb_long)

# A tibble: 3 x 6
  id      long   lat   elev name           wmo_id
  <chr>    <dbl> <dbl> <dbl> <chr>        <dbl>
1 ASN00086038  145. -37.7  78.4 essendon airport  95866
2 ASN00086077  145. -38.0  12.1 moorabbin airport  94870
3 ASN00086282  145. -37.7 113.  melbourne airport  94866
```

## 2.2. Creation and coercion

In many cases, spatio-temporal data arrive in separate tables for analysis. For example, in climate data, analysts may initially receive station data containing geographic location information, recorded variables and their recording periods. They can then query the temporal variables using the stations of interest to obtain the relevant temporal data. Alternatively, analyses may begin as purely spatial or temporal, and analysts may obtain additional temporal or spatial data to expand the result to spatio-temporal.

The function `make_cubble()` composes a `cubble` object from a spatial table (`spatial`) and a temporal table (`temporal`), along with three attributes introduced in the subsection 2.1: `key`, `index`, and `coords`. The following code creates the spatial `cubble` object shown previously:

```
R> make_cubble(spatial = stations, temporal = meteo,
+                 key = id, index = date, coords = c(long, lat))
```

Spatio-temporal data in other foreign objects can be coerced into a `cubble` object with the function `as_cubble()`. This includes casting from a `tibble` or `data.frame` with both spatial and temporal information, a NetCDF object, a `stars` object (Pebesma 2021), and a `sftime` object (Teickner *et al.* 2022). The two examples below show the casting from a tibble and a NetCDF object. The dataset `climate_flat` combines the spatial data, `stations`, with the temporal data, `meteo`, into a single tibble object. It can be coerced into a cubble using:

```
R> climate_flat |> as_cubble(key = id, index = date, coords = c(long, lat))

# cubble: key: id [3], index: date, nested form
# spatial: [144.8321, -37.98, 145.0964, -37.6655], Missing CRS!
# temporal: date [date], prcp [dbl], tmax [dbl], tmin [dbl]
#             id      long   lat elev name          wmo_id ts
#             <chr>    <dbl> <dbl> <dbl> <chr>        <dbl> <list>
# 1 ASN00086038 145. -37.7  78.4 essendon airport  95866 <tibble [10 x 4]>
# 2 ASN00086077 145. -38.0  12.1 moorabbin airport 94870 <tibble [10 x 4]>
# 3 ASN00086282 145. -37.7 113. melbourne airport 94866 <tibble [10 x 4]>
```

### 2.3. Functions and methods

Table 1 summarises the functions implemented in the `cubble` package and Table 2 details the methods implemented for each of the three cubble classes. The `cubble_df` class handles methods that behave consistently in both spatial (`spatial_cubble_df`) and temporal cubble (`temporal_cubble_df`). When the method has distinct behavior for the spatial cubble and temporal cubble, it is implemented separately for each subclass.

Table 1: Functions implemented in the `cubble` package

Category	Functions
base R	<code>[</code> , <code>[[&lt;-</code> , <code>names&lt;-</code>
tidyverse	<code>dplyr_row_slice</code> , <code>dplyr_col_modify</code> , <code>dplyr_reconstruct</code> , <code>select</code> , <code>mutate</code> , <code>arrange</code> , <code>filter</code> , <code>group_by</code> , <code>ungroup</code> , <code>summarise</code> , <code>select</code> , <code>slice</code> , <code>rowwise</code> , <code>rename</code> , <code>bind_rows</code> , <code>bind_cols</code> , <code>relocate</code> , <code>type_sum</code> , the slice family ( <code>slice_head</code> , <code>slice_tail</code> , <code>slice_max</code> , <code>slice_min</code> , <code>slice_sample</code> ) and the join family ( <code>left_join</code> , <code>right_join</code> , <code>inner_join</code> , <code>full_join</code> , <code>anti_join</code> , <code>semi_join</code> )
cubble	<code>as_cubble</code> , <code>cubble</code> , <code>make_cubble</code> , <code>check_key</code> face_temporal, <code>face_spatial</code> , <code>unfold</code> , <code>key</code> , <code>key_vars</code> , <code>key_data</code> , <code>index</code> , <code>index_var</code> , <code>coords</code> , <code>spatial</code> , <code>match_sites</code> , <code>match_spatial</code> , <code>match_temporal</code> , <code>geom_glyph</code> , <code>geom_glyph_box</code> , <code>geom_glyph_line</code> , <code>make_spatial_sf</code> , <code>make_temporal_tsibble</code> , <code>fill_gaps</code> , and <code>scan_gaps</code>

Table 2: Methods implemented for the three **cubble** classes.

Class	Methods
<code>cubble_df</code>	<code>[[&lt;-</code> , <code>dplyr_col_modify</code> , <code>key_data</code> , <code>key_vars</code> , <code>key</code> , <code>print</code>
<code>spatial_cubble_df</code>	<code>[, names&lt;-</code> , <code>tbl_sum</code> , <code>dplyr_reconstruct</code> ,
	<code>dplyr_row_slice</code> , <code>face_spatial</code> , <code>face_temporal</code> , <code>unfold</code> ,
	<code>arrange</code> , <code>rename</code> , <code>rowwise</code> , <code>group_by</code> , <code>ungroup</code> , <code>select</code> ,
	<code>spatial</code> , <code>summarise</code> , <code>unfold</code> , <code>update_cubble</code>
<code>temporal_cubble_df</code>	<code>[, names&lt;-</code> , <code>tbl_sum</code> , <code>arrange</code> , <code>dplyr_reconstruct</code> ,
	<code>dplyr_row_slice</code> , <code>face_spatial</code> , <code>face_temporal</code> , <code>unfold</code> ,
	<code>fill_gaps</code> , <code>group_by</code> , <code>ungroup</code> , <code>rename</code> , <code>rowwise</code> ,
	<code>scan_gaps</code> , <code>select</code> , <code>spatial</code> , <code>summarise</code> , <code>tbl_sum</code> ,
	<code>bind_rows</code> , <code>bind_cols</code> , <code>update_cubble</code>

Specifically, the pair of verbs, `face_temporal()` and `face_spatial()`, pivot the cubble object between the spatial and temporal face of the multivariate spatio-temporal cube, as illustrated in Figure 1. The code below uses the function `face_temporal()` and `face_spatial()` to transform between a spatial cubble (`cb_nested`) and a temporal one (`cb_long`), introduced earlier in subsection 2.1:

```
R> identical(face_temporal(cb_nested), cb_long)
```

```
[1] TRUE
```

```
R> identical(face_spatial(cb_long), cb_nested)
```

```
[1] TRUE
```

The pair of verbs are exact inverse and apply both functions on a cubble object will result in the object itself:

```
R> identical(face_spatial(face_temporal(cb_nested)), cb_nested)
```

```
[1] TRUE
```

```
R> identical(face_temporal(face_spatial(cb_long)), cb_long)
```

```
[1] TRUE
```

## 2.4. Compatibility with `tsibble` and `sf`

Analysts often have their preferred spatial or temporal data structure that they prefer to use for spatio-temporal analysis. For example, the `tbl_ts` class from the `tsibble` package (Wang *et al.* 2020) is commonly used in time series forecasting and similarly, the `sf` class (Pebesma 2018) is frequently used in spatial data science. In `cubble`, analysts have the flexibility to




Figure 1: An illustration of the function `face_temporal()` and `face_spatial()`: `face_temporal()` converts a spatial cubble (nested form) into a temporal cubble (long form) to focus on the temporal variables. Conversely, `face_spatial()` transforms a temporal cubble into a spatial one to emphasize spatial variables.

combine these two structures together by allowing the spatial component to also be an `sf` object and the temporal component to also be a `tsibble` object.

#### *Using a tsibble object as the temporal component*

The `key` and `index` arguments in a `cubble` object corresponds to the `tsibble` counterparts and they can be safely omitted, if the temporal component is a `tsibble` object, i.e. `meteo_ts` in the example below. The `tsibble` class from the input will be carried over to the `cubble` object:

```
R> ts_nested <- make_cubble(
+   spatial = stations, temporal = meteo_ts, coords = c(long, lat))
R> (ts_long <- face_temporal(ts_nested))

# cubble: key: id [3], index: date, long form, [tsibble]
# temporal: 2020-01-01 -- 2020-01-10 [1D], no gaps
# spatial: long [dbl], lat [dbl], elev [dbl], name [chr], wmo_id [dbl]
  id      date      prcp  tmax tmin
  <chr>    <date>    <dbl> <dbl> <dbl>
1 ASN00086038 2020-01-01     0  26.8  11
2 ASN00086038 2020-01-02     0  26.3  12.2
3 ASN00086038 2020-01-03     0  34.5  12.7
4 ASN00086038 2020-01-04     0  29.3  18.8
5 ASN00086038 2020-01-05    18  16.1  12.5
# i 25 more rows

R> class(ts_long)

[1] "temporal_cubble_df" "cubble_df"          "tbl_ts"
[4] "tbl_df"              "tbl"                "data.frame"
```

## 8 **cubble**: An R Package for Organizing and Wrangling Multivariate Spatio-temporal Data

The temporal cubble shows `[tsibble]` in the header to indicate the object also being in a `tbl_ts` class (`tsibble`). Methods applies to the `tbl_ts` class can also be applied to the temporal cubble objects, for example, checking whether the data contain temporal gaps:

```
R> ts_long %>% has_gaps()
```

```
# A tibble: 3 x 2
  id      .gaps
  <chr>   <lgl>
1 ASN00086038 FALSE
2 ASN00086077 FALSE
3 ASN00086282 FALSE
```

An existing `cubble` object can promote its temporal component to a `tsibble` object by applying `make_temporal_tsibble()`. The following code illustrates this with the object `cb_long` created in section 2.2 and the promoted cubble object is equal to the cubble object originally created from a `tsibble` object:

```
R> ts_long2 <- cb_long %>% make_temporal_tsibble()
R> identical(ts_long2, ts_long)
```

```
[1] TRUE
```

### *Using an sf object as the spatial component*

Similarly, an `sf` object can be supplied as the spatial component to create a `cubble` object, with the `coords` argument being omitted. This opens up the possibility to represent fixed area with polygons or multipolygons (see Applications 4.1) and the `coords` argument will be calculated as the centroids of the (multi)polygons. The `[sf]` print in the cubble header suggest an spatial component being also a `sf` object:

```
R> (sf_nested <- make_cubble(
+   spatial = stations_sf, temporal = meteo,
+   key = id, index = date))

# cubble:  key: id [3], index: date, nested form, [sf]
# spatial:  [144.8321, -37.98, 145.0964, -37.6655], WGS 84
# temporal: date [date], prcp [dbl], tmax [dbl], tmin [dbl]
  id      elev name  wmo_id  long     lat      geometry ts
  <chr>   <dbl> <chr> <dbl> <dbl> <dbl>      <POINT [°]> <list>
1 ASN00086038  78.4 essen~  95866  145. -37.7 (144.9066 -37.7276) <tibble>
2 ASN00086077  12.1 moora~  94870  145. -38.0  (145.0964 -37.98)  <tibble>
3 ASN00086282 113. melbo~  94866  145. -37.7  (144.8321 -37.6655) <tibble>

R> class(sf_nested)
```

```
[1] "spatial_cubble_df" "cubble_df"           "sf"
[4] "tbl_df"             "tbl"                 "data.frame"
```

The `sf` functions applicable to a `cubble` object have been listed in Table 1 and the following code shows how to perform coordinate transformation with `st_transform` on a cubble object:

```
R> sf_nested %>% sf::st_transform(crs = "EPSG:3857")

# cubble: key: id [3], index: date, nested form, [sf]
# spatial: [16122635.6225205, -4576600.8687746, 16152057.3639371,
#   -4532279.35567565], WGS 84
# temporal: date [date], prcp [dbl], tmax [dbl], tmin [dbl]
#           id      elev name  wmo_id long   lat      geometry ts
#           <chr>  <dbl> <chr> <dbl> <dbl> <dbl> <POINT [°]> <list>
1 ASN00086038  78.4 essen~  95866  145. -37.7 (16130929 -4541016) <tibble>
2 ASN00086077  12.1 moora~  94870  145. -38.0 (16152057 -4576601) <tibble>
3 ASN00086282  113. melbo~  94866  145. -37.7 (16122636 -4532279) <tibble>
```

The counterpart to promote the spatial component in an existing `cubble` to be an `sf` object is `make_spatial_sf()`:

```
R> cb_nested %>% make_spatial_sf()

# cubble: key: id [3], index: date, nested form, [sf]
# spatial: [144.8321, -37.98, 145.0964, -37.6655], WGS 84
# temporal: date [date], prcp [dbl], tmax [dbl], tmin [dbl]
#           id      long   lat elev name  wmo_id ts      geometry
#           <chr>  <dbl> <dbl> <dbl> <chr> <dbl> <list> <POINT [°]>
1 ASN00086038  145. -37.7  78.4 essen~  95866 <tibble> (144.9066 -37.7276)
2 ASN00086077  145. -38.0  12.1 moora~  94870 <tibble> (145.0964 -37.98)
3 ASN00086282  145. -37.7  113. melbo~  94866 <tibble> (144.8321 -37.6655)
```

## 2.5. Comparison to other spatio-temporal classes

In R, there are other existing spatio-temporal data structure and this section compares and contrasts `cubble` with other existing alternative, specifically `stars` and `sftime`. The `stars` package (Pebesma 2021) uses an array structure, as oppose to tibble, to represent multivariate spatio-temporal data. While both `stars` and `cubble` support vector and raster data, it is a matter of choice on which structure to use given the application. Analysts working on satellite imageries may prefer the array structure in `stars`, while others originally working with spatio-temporal data in 2D data frames may find `cubble` easier to adopt from their existing computing workflow.

The `sftime` package (Teickner *et al.* 2022) also builds from a tibble object and its focus is on handling irregular spatio-temporal data. This means `sftime` can also handle full space-time grids and sparse space-time layouts represented in `cubble`, but `cubble` uses nesting to avoid

storing spatial variables repetitively at each timestamp. This provides memory efficiency when data is observed frequent, i.e. daily or sub-daily, or the spatial geometry is expensive to repeat, i.e. polygons or multipolygons. Consider the `climate_aus` data in the **cubble** package with 639 stations observed daily in a single year 2020. The created `sftime` object is approximately 14 times larger than the corresponding `cubble` object (118.24 MB vs. 8.52 MB).

### 3. Other features and considerations

#### 3.1. Data fusion and matching

One common task when working with spatio-temporal data is to match nearby sites. For example, we may want to verify the location of an old list of stations with current stations, or we may want to match the data from different data sources. In cubble, spatial and temporal matching are performed using the functions `match_spatial()` and `match_temporal()`. The `match_spatial()` function calculates the spatial distance between observations in two cubble objects. Various distance measures are available (check `sf::st_distance`). Analysts can specify the number of matched groups to output using the `spatial_n_group` argument (default to 4 groups) and the number of matches per group using the `spatial_n_group` argument (default to 1, one-to-one matching). The syntax to use `match_spatial()` is:

```
match_spatial(<cubble_obj1>, <cubble_obj2>, ...)
```

The function `match_temporal()` calculates the similarity between time series within spatially matched groups. Two identifiers are required: one for separating each matched group (`match_id`) and one for separating the two data sources (`data_id`). The argument `temporal_by` uses the `by` syntax from dplyr's `*_join` to specify the temporal variables to match.

The similarity score between two time series is calculated using a matching function, which can be customised by the analysts. The matching function takes two time series as a list and returns a single numerical score. This allows for flexibility in using existing time series distance calculation implementation. By default, cubble implements a simple peak matching algorithm (`match_peak`) that counts the number of peaks in two time series that fall within a specified temporal window. The syntax to use `match_temporal()` is

```
match_temporal(
  <cubble_obj_from_match_spatial>,
  data_id = , match_id = ,
  temporal_by = c("..." = "...")
)
```

#### 3.2. Interactive graphics

The workflow with the `cubble` class works well with an interactive graphics pipeline (e.g., Buja *et al.* (1988); Buja *et al.* (1996); Sutherland *et al.* (2000); Xie *et al.* (2014); Cheng *et al.* (2016)) that is available in R with the package `crosstalk` (Cheng and Sievert 2021).




Figure 2: Linking between multiple plots. The line plots and the map are constructed from shared `crosstalk` objects. When a station is selected on the map (a), the corresponding row in the spatial `cubicle` will be activated. This will link to all the rows with the same id in the temporal `cubicle` (b) and update the line plot (c).

Figure 2 illustrates the linking between a map and multiple time series in a `cubicle` object. Both the map, generated from the spatial `cubicle`, and the time series plot, generated from the temporal `cubicle`, are shared `crosstalk` objects. When a user selects a location on the map (a), the corresponding site is highlighted. This selection activates a row in the spatial `cubicle`, which is then communicated to the temporal `cubicle`, resulting in the selection of all observations with the same ID in (b). Consequently the temporal `cubicle` highlights the corresponding series in the time series plot (c). Linking can also be initiated from the time series plot, by selecting points on the time series. This action activates rows with the same ID in the temporal `cubicle` and the row in the spatial `cubicle`, which will be highlighted on the map. Please refer to the appendix for the corresponding illustration.

### 3.3. Spatio-temporal transformations

Sometimes, we wish to communicate spatial and temporal information collectively through visualisation. This can be achieved through several graphical displays: one can make faceted maps across time, creating map animations, or constructing interactive graphics to link between map and time series plot. The concept of *glyph maps* was initially proposed in Wickham *et al.* (2012). The underlying idea is to transform the temporal coordinates into spatial coordinates so that time series plot can be displayed on the map.

The package `GGally` initially implemented the *glyph map*. It uses `glyphs()` to calculate the axis transformation and then uses `geom_polygon()` to draw the map. In `cubicle`, a ggproto implementation `geom_glyph()` performs the linear algebra internally as data transformation. The `geom_glyph()` requires four aesthetics: `x_major`, `y_major`, `x_minor`, and `y_minor`. The major axes are the outer spatial coordinates and the minor axes are the inner/ temporal

coordinates:

```
data |>
  ggplot() +
  geom_glyph(aes(x_major = ..., x_minor = ...,
                  y_major = ..., y_minor = ...))
```

Some useful controls over the glyph map includes:

- polar coordinate glyph maps with `polar = TRUE`,
- adjust glyph size with arguments `width` and `height`,
- glyph scale to fixed (`global_rescale` defaults to `TRUE`) or free, and
- reference boxes and lines with `geom_glyph_box()` and `geom_glyph_line()`.

## 4. Applications

Five examples are chosen to illustrate different aspects of the **cubble** package: creating a **cubble** object from two Coronavirus (COVID) data tables with the complication of differing location names, using spatial transformations to make a glyph map of seasonal temperature changes, matching river level data and weather station records for analysis of water supply, reading NetCDF format data to reproduce a climate reanalysis plot, and the workflow to create complex interactive linked plots.

### 4.1. Victoria COVID spatio-temporal incidence and spread

Since the start of the pandemic, the Victoria State Government in Australia has been providing daily COVID counts per Local Government Area (LGA). This data can be combined with map polygon data, available from the Australian Bureau of Statistics (ABS), to visualize COVID incidence and spread. In the **cubble** package, the COVID count data (`covid`) and the LGA information (`lga`) are available as a `tsibble` object and an `sf` object respectively.

A **cubble** object can be created from separate spatial and temporal component using the `make_cubble()` function, introduced in Section 2.2. The `by` argument is used to specify the joining variable from the two component using the `*_join()` by syntax:

```
R> cb <- make_cubble(lga, covid, by = c("lga_name_2018" = "lga"))
```

Warning: `st_centroid` assumes attributes are constant over geometries

! Some sites in the spatial table don't have temporal information

! Some sites in the temporal table don't have spatial information

! Use `check\_key()` to check on the unmatched key

The **cubble** is created only with sites having both spatial and temporal information

The warning message suggests the slight difference of LGA encoding used by Victoria government and ABS and prompts analysts to use the function `check_key()` to identify the mismatches:

```
R> (check_res <- check_key(
+   spatial = lga, temporal = covid,
+   by = c("lga_name_2018" = "lga")
+ ))
```

\$paired

```
# A tibble: 78 x 2
  spatial      temporal
  <chr>        <chr>
  1 Alpine (S)  Alpine (S)
  2 Ararat (RC) Ararat (RC)
  3 Ballarat (C) Ballarat (C)
  4 Banyule (C) Banyule (C)
  5 Bass Coast (S) Bass Coast (S)
# i 73 more rows
```

\$potential\_pairs

```
# A tibble: 2 x 2
  spatial      temporal
  <chr>        <chr>
  1 Kingston (C) (Vic.) Kingston (C)
  2 Latrobe (C) (Vic.) Latrobe (C)
```

\$others

```
$others$spatial
[1] "No usual address (Vic.)"
[2] "Migratory - Offshore - Shipping (Vic.)"
```

\$others\$temporal

```
[1] "Interstate" "Overseas"    "Unknown"
```

The result of the `check_key()` function is a list containing three elements: 1) matched keys from both tables, 2) potentially paired keys, and 3) others. Here, the main mismatch arises from the two LGAs: Kingston and Latrobe (Kingston is a LGA in both Victoria and South Australia and Latrobe is a LGA in both Victoria and Tasmania). Analysts can modify the input spatial and temporal data accordingly and recreate the cubble object:

```
R> lga2 <- lga />
+   rename(lga = lga_name_2018) />
+   mutate(lga = ifelse(lga == "Kingston (C) (Vic.)", "Kingston (C)", lga),
+         lga = ifelse(lga == "Latrobe (C) (Vic.)", "Latrobe (C)", lga)) />
+   filter(!lga %in% check_res$others$spatial)
R>
```

```
R> covid2 <- covid |> filter(!lga %in% check_res$others$temporal)
R>
R> (cb <- make_cubble(spatial = lga2, temporal = covid2))

# cubble:   key: lga [80], index: date, nested form, [sf]
# spatial:  [140.961682, -39.1591895, 149.976291, -33.9804256], WGS 84
# temporal: date [date], n [dbl], avg_7day [dbl]
# i 75 more rows
```

lga	long	lat	geometry	ts
<chr>	<dbl>	<dbl>	<MULTIPOLYGON [°]>	<list>
1 Alpine (S)	147.	-36.9	((146.7258 -36.45922, 146.8033 -36.45139~	<tbl_ts>
2 Ararat (RC)	143.	-37.5	((143.1807 -37.73152, 143.1793 -37.73242~	<tbl_ts>
3 Ballarat (C)	144.	-37.5	((143.6622 -37.57241, 143.6686 -37.53844~	<tbl_ts>
4 Banyule (C)	145.	-37.7	((145.1357 -37.74091, 145.1331 -37.74281~	<tbl_ts>
5 Bass Coast (S)	146.	-38.5	((145.5207 -38.30667, 145.5418 -38.30938~	<tbl_ts>

## 4.2. Australian historical maximum temperature

The Global Historical Climatology Network (GHCN) provides daily climate measures for stations worldwide. In the **cubble** package, the dataset `historical_tmax` contains daily maximum temperature data for 75 stations in Australia, covering two periods: 1971-1975 and 2016-2020. The objective of this example is to compare the changes in maximum temperature between these two periods using a glyph map.

To prevent overlapping of weather stations on the map, we subset the stations based on minimum distance of 50km between them. This can be done by first promoting the spatial cubble to also be an sf object with `make_spatial_sf()`, calculating the distance matrix using the sf function `st_distance()`, and filtering the stations with dplyr's `filter()`:

```
R> a <- historical_tmax %>% make_spatial_sf() %>% st_distance()
R> a[upper.tri(a, diag = TRUE)] <- 1e6
R>
R> (tmax <- historical_tmax |>
+   filter(rowSums(a < units::as_units(50, "km")) == 0))

# cubble:   key: id [54], index: date, nested form
# spatial:  [141.2652, -39.1297, 153.3633, -28.9786], Missing CRS!
# temporal: date [date], tmax [dbl]
# i 49 more rows
```

id	long	lat	elev	name	wmo_id	ts
<chr>	<dbl>	<dbl>	<dbl>	<chr>	<dbl>	<list>
1 ASN00047016	141.	-34.0	43	lake victoria storage	94692	<tibble>
2 ASN00047019	142.	-32.4	61	menindee post office	94694	<tibble>
3 ASN00048015	147.	-30.0	115	bewarrina hospital	95512	<tibble>
4 ASN00048027	146.	-31.5	260	cobar mo	94711	<tibble>
5 ASN00048031	149.	-29.5	145	collarenebri (albert st)	95520	<tibble>

Next, daily maximum temperature is averaged into monthly series for each periods in the temporal cubble. The last step with `unfold()` moves the two coordinate columns `long`, `lat` into the temporal cubble, preparing the data for the glyph map:

```
R> (tmax <- tmax |>
+   face_temporal() |>
+   group_by(
+     yearmonth = tsibble::make_yearmonth(
+       year = ifelse(lubridate::year(date) > 2015, 2016, 1971),
+       month = lubridate::month(date))
+   ) |>
+   summarise(tmax = mean(tmax, na.rm = TRUE)) %>%
+   mutate(group = as.factor(lubridate::year(yearmonth)),
+         month = lubridate::month(yearmonth)) %>%
+   unfold(long, lat))

# cubble:  key: id [54], index: yearmonth, long form
# temporal: 1971 Jan -- 2016 Dec [1M], has gaps!
# spatial:  long [dbl], lat [dbl], elev [dbl], name [chr], wmo_id [dbl]
  yearmonth id      tmax group month long  lat
    <mth> <chr>    <dbl> <fct> <dbl> <dbl> <dbl>
 1 1971 Jan ASN00047016 31.1 1971      1 141. -34.0
 2 1971 Jan ASN00047019 33.1 1971      1 142. -32.4
 3 1971 Jan ASN00048015 33.9 1971      1 147. -30.0
 4 1971 Jan ASN00048027 32.5 1971      1 146. -31.5
 5 1971 Jan ASN00048031 33.3 1971      1 149. -29.5
# i 1,276 more rows
```

A quick check on the number of observations for each location is made, revealing that there are several with less than 24 observations – these stations lack temperature values for some months. In this example, those stations are removed by switching to the spatial `cubble` to operate on the spatial component over time, and then, move back into the temporal `cubble` (to make the glyph map):

```
R> tmax <- tmax |>
+   face_spatial() |>
+   rowwise() |>
+   filter(nrow(ts) == 24) |>
+   face_temporal()
```

The following code creates the glyph map (a) in Figure ?? with additional ggplot2 code for highlighting the single station, Cobar and styling:

```
R> nsw_vic <- ozmaps::abs_stc |>
+   filter(NAME %in% c("Victoria", "New South Wales"))
R>
```




Figure 3: Comparison of Average maximum temperature between 1971-1975 and 2016-2020 for 54 stations in Victoria and New South Wales, Australia. (a) and (b): the monthly temperature series for the two periods in a glyph map and for a single station Cobar, highlighted in orange in the glyph map. (c) and (d): the difference series between the two periods (2016s minus 1971s) in a glyph map and for station Cobar. The grey horizontal line marks zero difference. The glyph map displaying the difference series (c) reveals more pronounced changes between the two periods, with many inland locations in New South Wales show an increased temperature in late summer (Jan-Feb) in recent years.

```
R> tmax %>%
+   ggplot(aes(x_major = long, x_minor = month,
+               y_major = lat, y_minor = tmax,
+               group = interaction(id, group))) +
+   geom_sf(data = nsw_vic, ..., inherit.aes = FALSE) +
+   geom_glyph_box(width = 0.8, height = 0.3) +
+   geom_glyph(aes(color = group), width = 0.8, height = 0.3) +
+   ...
```

### 4.3. River levels and rainfall in Victoria

One common task when working with spatio-temporal data is to match nearby sites. For example, we may want to verify the location of an old list of stations with current stations, or we may want to match the data from different data sources. In this example, we will

introduce the spatial and temporal matching in cubble using an example on matching river level data with precipitation in Victoria, Australia.

The water level data collected by the Bureau of Meteorology, can be compared with the precipitation since rainfall can directly impact water level in river. Both `climate_vic` and `river` are cubble objects, and we can obtain a summary of the 10 closest pairs between them:

```
R> res_sp <- match_spatial(climate_vic, river, spatial_n_group = 10)
R> print(res_sp, n = 20)

# A tibble: 10 x 4
  from      to    dist group
  <chr>     <chr>   [m] <int>
1 ASN00088051 406213 1838.     1
2 ASN00084145 222201 2185.     2
3 ASN00085072 226027 3282.     3
4 ASN00080015 406704 4034.     4
5 ASN00085298 226027 4207.     5
6 ASN00082042 405234 6153.     6
7 ASN00086038 230200 6167.     7
8 ASN00086282 230200 6928.     8
9 ASN00085279 224217 7431.     9
10 ASN00080091 406756 7460.    10
```

The result can also be returned as cubble objects by setting the argument `return_cubble = TRUE`. The output is be a list where each element is a paired cubble object. To combine all the results into a single cubble, you can use `bind_rows()`. In the case when a site in the second cubble (the `river` data here) is matched to two stations in the first cubble (`climate_vic` here), the binding may not be successful since cubble requires unique rows in the nested form. In the summary table above, the river station 226027 is matched to more than one weather station: ASN00085072 (group 3) and ASN00085298 (group 5). Similarly, the river station 230200 is matched in group 7 and 8). In such cases, you can either deselect one pair before combining, or work with the list output with the `purrr::map` syntax:

```
R> res_sp <- match_spatial(
+   climate_vic, river,
+   spatial_n_group = 10, return_cubble = TRUE)
R> str(res_sp, max.level = 0)

List of 10

R> (res_sp <- res_sp[-c(5, 8)] %>% bind_rows())

# cubble:  key: id [16], index: date, nested form, [sf]
# spatial:  [144.5203, -38.144913, 148.4667, -36.128657], WGS 84
# temporal: date [date], prcp [dbl], tmax [dbl], tmin [dbl]
#             id      long     lat   elev name   wmo_id ts       type      geometry
```

```

<chr> <dbl> <dbl> <dbl> <chr> <dbl> <list>   <chr>           <POINT [°]>
1 ASN00~ 145. -37.0 290  rede~ 94859 <tibble> clim~ (144.5203 -37.0194)
2 406213 145. -37.0 NA   CAMP~      NA <tibble> river (144.5403 -37.01512)
3 ASN00~ 148. -37.7 62.7 orbo~ 95918 <tibble> clim~ (148.4667 -37.6922)
4 222201 148. -37.7 NA   SNOW~      NA <tibble> river (148.451 -37.70739)
5 ASN00~ 147. -38.1  4.6 east~ 94907 <tibble> clim~ (147.1322 -38.1156)
# i 11 more rows
# i 2 more variables: group <int>, dist [m]

```

For temporal matching, we match the variable `Water_course_level` from the river data to `prcp` in the weather station data. The variable `group` and `types` identify the matching group and the two datasets:

```

R> (res_tm <- res_sp %>%
+   match_temporal(
+     data_id = type, match_id = group,
+     temporal_by = c("prcp" = "Water_course_level")))

# A tibble: 8 x 2
  group match_res
  <int>    <dbl>
1     1        30
2     2         5
3     3        14
4     4        20
5     6        23
# i 3 more rows

```

Similarly, the cubble output can be returned using the argument `return_cubble = TRUE`. Here we select the four pairs with the highest number of matching peaks:

```

R> res_tm <- res_sp %>%
+   match_temporal(
+     data_id = type, match_id = group,
+     temporal_by = c("prcp" = "Water_course_level"),
+     return_cubble = TRUE)
R> (res_tm <- res_tm %>% bind_rows() %>% filter(group %in% c(1, 7, 6, 9)))

# cubble: key: id [8], index: date, nested form, [sf]
# spatial: [144.5203, -37.8817, 147.572223, -36.8472], WGS 84
# temporal: date [date], matched [dbl]
# # # 
  id      long   lat elev name  wmo_id type           geometry group
  <chr>    <dbl> <dbl> <dbl> <chr> <dbl> <chr>           <POINT [°]> <int>
1 ASN00088~ 145. -37.0 290  rede~ 94859 clim~ (144.5203 -37.0194)    1
2 406213    145. -37.0 NA   CAMP~      NA river (144.5403 -37.01512)    1
3 ASN00082~ 146. -36.8 502  stra~ 95843 clim~ (145.7308 -36.8472)    6

```

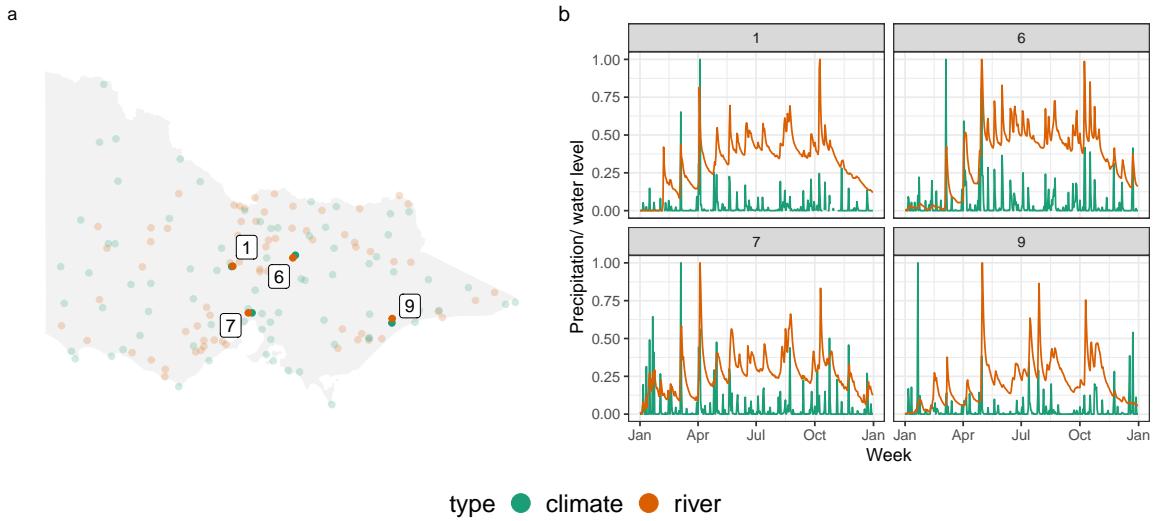


Figure 4: Weather stations and river gauges with matched pairs labelled on the map (a) and plotted across time (b). Precipitation and water level have been standardised between 0 and 1 to be displayed on the same scale. The water level reflects the increase in precipitation. The numbers (1, 7, 6, 9) indicate the group index derived from spatial matching, only those that were selected by temporal matching are shown here.

```

4 405234      146. -36.9  NA   SEVE~      NA river (145.6828 -36.88701)      6
5 ASN00086~    145. -37.7  78.4 esse~  95866 clim~ (144.9066 -37.7276)      7
# i 3 more rows
# i 3 more variables: dist [m], ts <list>, match_res <dbl>

```

Figure 4 shows four matched pairs on the map (a) and standardized data as time series (b) with concurrent increasing.

#### 4.4. ERA5: climate reanalysis data

The ERA5 reanalysis (Hersbach *et al.* 2020) provides hourly estimates of atmospheric, land and oceanic climate variables on a global scale. The hourly pressure level data in the NetCDF format can be downloaded from Copernicus Climate Data Store (CDS) (Copernicus Climate Change Service 2022) or via the **ecmwfr** package (Hufkens *et al.* 2019). This example reproduces the row created from ERA5 reanalysis data shown in Figure 19 by Hersbach *et al.* (2020). The plot shows the southern polar vortex splitting into two on 2002-09-26, and further splitting into four on 2002-10-04. Further explanation of why this is interesting can be found in the figure source, and also in Simmons *et al.* (2020) and Simmons *et al.* (2005).

To work with NetCDF data in R, several packages are available, including **ncdf4**, **RNetCDF**, and **tidync**. The following code converts a NetCDF object of class **ncdf4** into a cubicle object:

```
R> raw <- ncdf4::nc_open(here::here("data/era5-pressure.nc"))
```

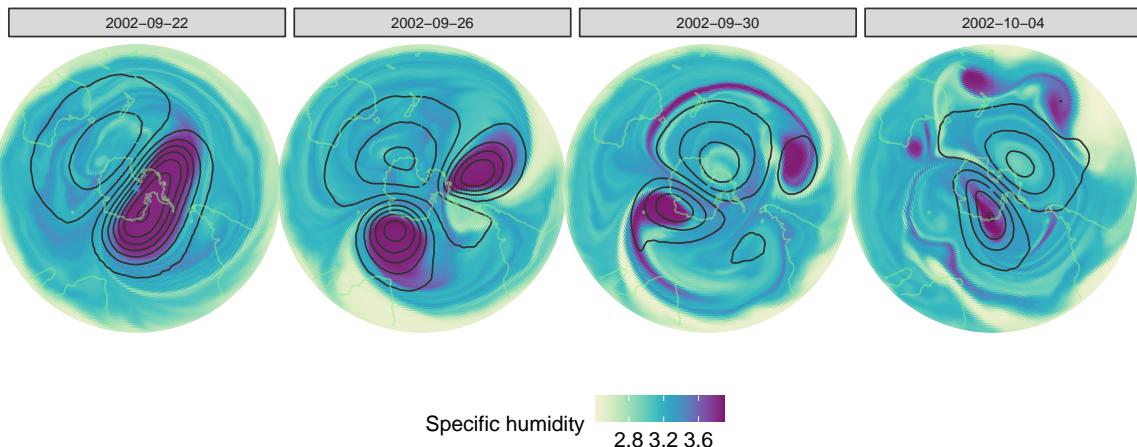


Figure 5: A reproduction of the second row (ERA5 data) of Figure 19 in Hersbach et al (2020) to illustrate the break-up of southern polar vortex in late September and early October 2002. The polar vortex, signalled by the high specific humidity, splits into two on 2002-09-26 and further splits into four on 2002-10-04.

Analysts can extract a subset of the NetCDF data with the arguments `vars`, `long_range` and `lat_range`. In this example, the variables `q` (specific humidity) and `z` (geopotential) are selected and the coordinates are subsetted to every degree in longitude and latitude:

```
R> (dt <- as_cubble(raw, vars = c("q", "z"),
+                     long_range = seq(-180, 180, 1), lat_range = seq(-88, -15, 1)))

# cubble:   key: id [26640], index: time, nested form
# spatial:  [-180, -88, 179, -15], Missing CRS!
# temporal: time [date], q [dbl], z [dbl]
      id long  lat ts
      <int> <dbl> <dbl> <list>
1     1 -180  -15 <tibble [8 x 3]>
2     2 -179  -15 <tibble [8 x 3]>
3     3 -178  -15 <tibble [8 x 3]>
4     4 -177  -15 <tibble [8 x 3]>
5     5 -176  -15 <tibble [8 x 3]>
# i 26,635 more rows
```

Once the NetCDF data is coerced into a cubble object, subsequent analysis can be conducted to filter on the date of interest, scale the variable specific humidity and create visualisation in ggplot as in Figure 5.

#### 4.5. Australian temperature range

Interactive graphics can be useful because they make it possible to look at the data in multiple ways on-the-fly. This is especially important for spatio-temporal data, where we would like to interactively connect spatial and temporal displays. This example describes the process of

using the **cubicle** package with the **crosstalk** package to build an interactive display connecting a map of Australia, with ribbon plots of temperature range observed at the stations. The purpose is to explore the variation of monthly temperature range over the country. Figure 6 shows three snapshots of the interactivity.

The key steps are to convert both the nested and long forms of the data into shared **crosstalk** objects, and to plot these side-by-side. The two are linked by the station identifier.

```
clean <- climate_full |> ...

nested <- clean |> SharedData$new(~id, group = "cubicle")
long <- face_temporal(clean) |> SharedData$new(~id, group = "cubicle")

p1 <- nested |> ...
p2 <- long |> ...

crosstalk::bscols(plotly::ggplotly(p1), plotly::ggplotly(p2), ...)
```

Plot (a) shows the initial state of the interactive display: all locations are shown as dots on the map, coloured by temperature range, and the right plot shows the ribbons representing maximum to minimum for all stations. In plot (b) the “Mount Elizabeth” station, which shows a high variance colour on the initial map, is selected on the map and this produces the ribbon on the right. In plot (c) the lowest temperature in August is selected, which is “Thredbo” station on the left map. It was surprising to us that this was not a station in Tasmania, so for comparison a station in Tasmania is selected on the map to show in relation to Thredbo. We can see that Thredbo has a bigger winter dip in temperature, and although Tasmania is cold generally, its temperatures are more constant

## 5. Conclusion

This paper presents the R package **cubicle** for organizing, wrangling and visualizing spatio-temporal data. The package introduces a new data class, **cubicle**, consisting of two subclasses, spatial cubicle and a temporal cubicle, to organise spatio-temporal data in two different formats within the tidy data framework. The data structure and functions introduced in this package can be used and combined with existing spatial and temporal data analysis packages such as **sf** and **tsibble**, data wrangling packages such as **dplyr**, and visualization packages such as **ggplot2**, **plotly**, and **leaflet**.

The paper includes numerous examples to illustrate the utility of **cubicle** as a data structure for spatio-temporal analysis. These examples cover different tasks of a typical analysis workflow: handling data with spatial and temporal misalignment, matching data from multiple sources, and creating both static and interactive spatio-temporal visualisation. For future directions, other commonly-used spatial or temporal data structures can be integrated into **cubicle** to extend analysts’ familiar spatial and temporal toolkit to spatio-temporal.

## 6. Acknowledgement

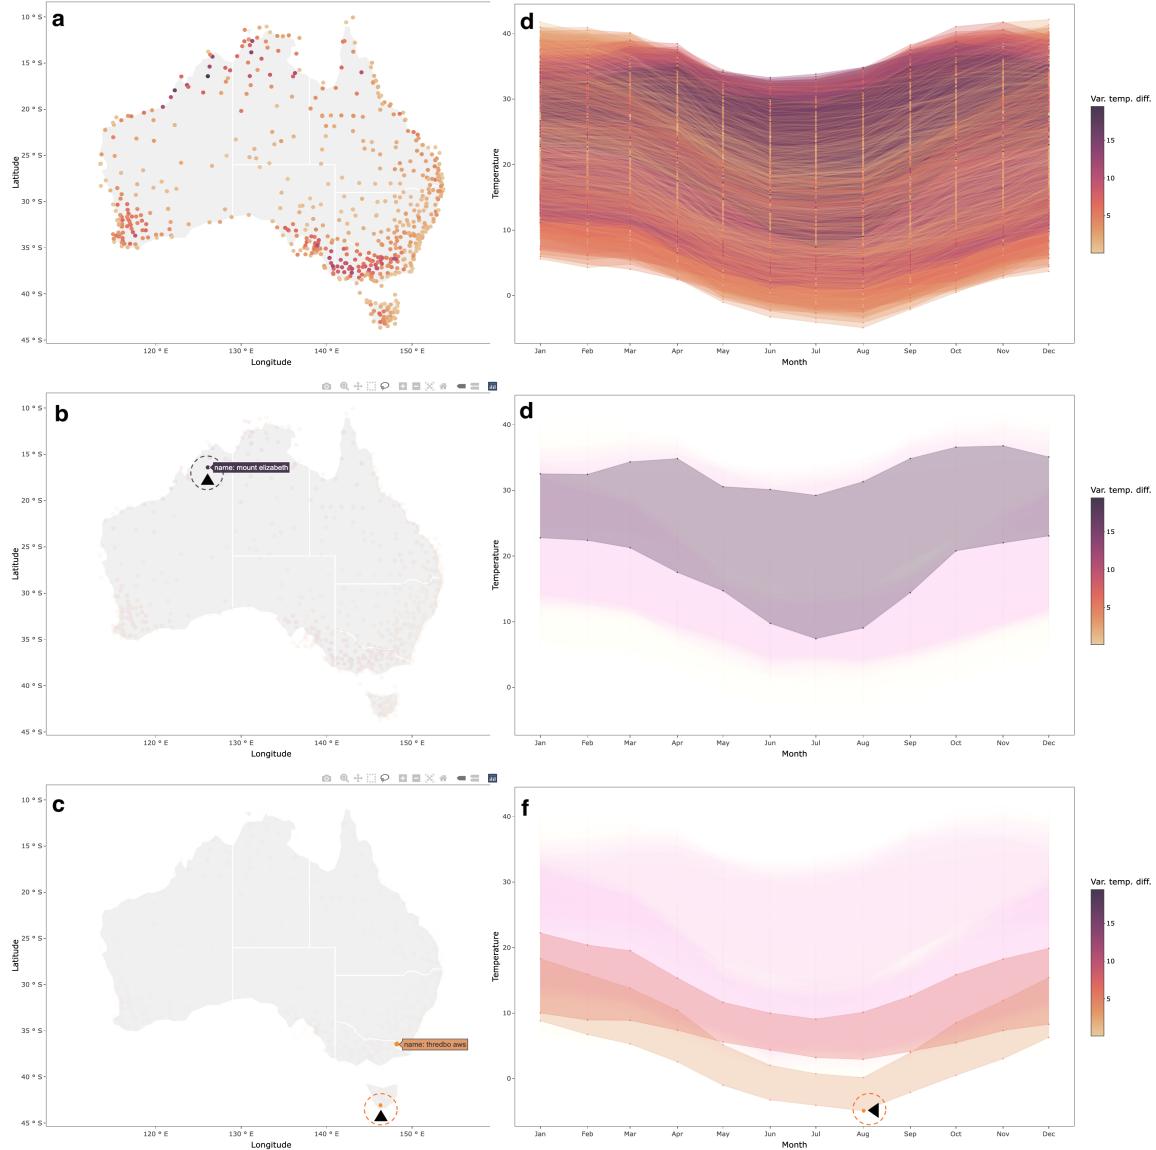


Figure 6: Exploring temperature variation using linking of a map and seasonal display. Each row is a screen dump of the process. The top row shows all locations and all temperature profiles. Selecting a particular location on the map (here Mount Elizabeth) produces the plot in the second row. The maximum and minimum temperatures are shown using a ribbon. The bottom row first selects the lowest temperature in August in the seasonal display, which highlights the corresponding station on the map (Thredbo). Another station, located in the Tasmania Island, is then selected to compare its temperature variation with the Thredbo station.

This work is funded by a Commonwealth Scientific and Industrial Research Organisation (CSIRO) Data61 Scholarship and started while Nicolas Langrené was affiliated with CSIRO's Data61. The article is created using the package **knitr** (Xie 2015) and **rmarkdown** (Xie *et al.* 2018) in R with the **rticles::jss\_article** template. The source code for reproducing this paper can be found at: <https://github.com/huizehang-sherry/paper-cubble>.

## References

- Buja A, Asimov D, Hurley C (1988). “Elements of A Viewing Pipeline.” *Dynamic Graphics Statistics*, p. 277.
- Buja A, Cook D, Swayne DF (1996). “Interactive High-dimensional Data Visualization.” *Journal of Computational and Graphical Statistics*, **5**(1), 78–99. URL <https://doi.org/10.2307/1390754>.
- Cheng J, Sievert C (2021). **crosstalk**: Inter-Widget Interactivity for HTML Widgets. R package version 1.1.1, URL <https://CRAN.R-project.org/package=crosstalk>.
- Cheng X, Cook D, Hofmann H (2016). “Enabling Interactivity on Displays of Multivariate Time Series and Longitudinal Data.” *Journal of Computational and Graphical Statistics*, **25**(4), 1057–1076. URL <https://doi.org/10.1080/10618600.2015.1105749>.
- Copernicus Climate Change Service (2022). “Climate Data Store.” Online; accessed 24 June 2022, URL <https://cds.climate.copernicus.eu/#!/home>.
- Hersbach H, Bell B, Berrisford P, Hirahara S, Horányi A, Muñoz-Sabater J, Nicolas J, Peubey C, Radu R, Schepers D, *et al.* (2020). “The ERA5 Global Reanalysis.” *Quarterly Journal of the Royal Meteorological Society*, **146**(730), 1999–2049.
- Hufkens K, Stauffer R, Campitelli E (2019). “The **ecwmfr** Package: An Interface to ECMWF API Endpoints.” URL <https://bluegreen-labs.github.io/ecmwfr/>.
- Pebesma E (2012). “**spacetime**: Spatio-Temporal Data in R.” *Journal of Statistical Software*, **51**(7), 1–30. URL <https://doi.org/10.18637/jss.v051.i07>.
- Pebesma E (2018). “Simple Features for R: Standardized Support for Spatial Vector Data.” *R Journal*, **10**(1), 439.
- Pebesma E (2021). **stars**: Spatiotemporal Arrays, Raster and Vector Data Cubes. R package version 0.5-2, URL <https://CRAN.R-project.org/package=stars>.
- Pebesma E, Bivand R (2022). “CRAN Task View: Handling and Analyzing Spatio-Temporal Data.” Version 2022-03-07, URL <https://CRAN.R-project.org/view=SpatioTemporal>.
- Simmons A, Hortal M, Kelly G, McNally A, Untch A, Uppala S (2005). “ECMWF Analyses and Forecasts of Stratospheric Winter Polar Vortex Breakup: September 2002 in the Southern Hemisphere and Related Events.” *Journal of the Atmospheric Sciences*, **62**(3), 668 – 689. doi:[10.1175/JAS-3322.1](https://doi.org/10.1175/JAS-3322.1). URL <https://journals.ametsoc.org/view/journals/atsc/62/3/jas-3322.1.xml>.

- Simmons A, Soci C, Nicolas J, Bell B, Berrisford P, Dragani R, Flemming J, Haimberger L, Healy S, Hersbach H, Horányi A, Inness A, Munoz-Sabater J, Radu R, Schepers D (2020). “Global Stratospheric Temperature Bias and Other Stratospheric Aspects of ERA5 and ERA5.1.” (859). doi:10.21957/rcxqfmg0. URL <https://www.ecmwf.int/node/19362>.
- Sutherland P, Rossini A, Lumley T, Lewin-Koh N, Dickerson J, Cox Z, Cook D (2000). “**Orca**: A Visualization Toolkit for High-dimensional Data.” *Journal of Computational and Graphical Statistics*, **9**(3), 509–529. URL <https://www.tandfonline.com/doi/abs/10.1080/10618600.2000.10474896>.
- Teickner H, Pebesma E, Graeler B (2022). *sftime: Classes and Methods for Simple Feature Objects that Have a Time Column*. <Https://r-spatial.github.io/sftime/>, <https://github.com/r-spatial/sftime>.
- Wang E, Cook D, Hyndman RJ (2020). “A New Tidy Data Structure to Support Exploration and Modeling of Temporal Data.” *Journal of Computational and Graphical Statistics*, **29**(3), 466–478. URL <https://doi.org/10.1080/10618600.2019.1695624>.
- Wickham H (2014). “Tidy Data.” *Journal of Statistical Software*, **59**(10), 1–23. URL <https://doi.org/10.18637/jss.v059.i10>.
- Wickham H, Hofmann H, Wickham C, Cook D (2012). “Glyph-Maps for Visually Exploring Temporal Patterns in Climate Data and Models.” *Environmetrics*, **23**(5), 382–393. doi:10.1002/env.2152.
- Xie Y (2015). *Dynamic Documents with R and knitr*. 2nd edition. Chapman and Hall/CRC, Boca Raton, Florida. ISBN 978-1498716963, URL <https://yihui.name/knitr/>.
- Xie Y, Allaire J, Grolemund G (2018). *R Markdown: The Definitive Guide*. Chapman and Hall/CRC, Boca Raton, Florida. ISBN 978-1138359338, URL <https://bookdown.org/yihui/rmarkdown>.
- Xie Y, Hofmann H, Cheng X (2014). “Reactive Programming for Interactive Graphics.” *Statistical Science*, **29**(2), 201 – 213. URL <https://doi.org/10.1214/14-STS477>.

## Affiliation:

H. Sherry Zhang  
Monash University  
21 Chancellors Walk, Clayton VIC 3800 Australia  
E-mail: [huize.zhang@monash.edu](mailto:huize.zhang@monash.edu)

Dianne Cook  
Monash University  
21 Chancellors Walk, Clayton VIC 3800 Australia  
E-mail: [dicook@monash.edu](mailto:dicook@monash.edu)

Ursula Laa  
University of Natural  
Resources and Life Sciences  
Gregor-Mendel-Straße 33, 1180 Wien, Austria  
E-mail: [ursula.laa@boku.ac.at](mailto:ursula.laa@boku.ac.at)

Nicolas Langrené  
BNU-HKBU  
United International College  
2000 Jintong Road, Tangjiawan, Zhuhai, Guangdong Province, China  
E-mail: [nicolaslangrene@uic.edu.cn](mailto:nicolaslangrene@uic.edu.cn)

Patricia Menéndez  
Monash University  
21 Chancellors Walk, Clayton VIC 3800 Australia  
E-mail: [patricia.menendez@monash.edu](mailto:patricia.menendez@monash.edu)